Sharko

Loan What-If Tool

Documentation and User Guide

What it does, how to run it, how to read the results, how steady and reliable they are, and how the code is organized.

MODESIM course project (SF241)
October 2026

Amounts in this document are in Indian Rupees (INR).

Contents

1. What is Sharko? 3

A worked example 3

What Sharko is not 3

2. How Sharko works 4

What stays fixed and what can change 4

The three ways of searching 4

How Sharko decides what is "closest" 4

3. Quick start 5

How long things take 6

4. Using the check command 8

The questions 8

Skipping the questions with flags 8

Reading the result in the terminal 9

5. Your report file 12

6. The experiment and the charts 13

The four charts and how to read them 13

Files the experiment and report write 15

7. How steady and how reliable are the results? 16

Negative assets and different random splits 16

Different kinds of applicant 16

Do the percentages mean what they say? 17

Sample outputs 18

8. Command reference 19

What the scores printed by train and evaluate mean 19

Exit codes 20

9. How the code is organized 20

File by file 20

Where do I change...? 21

Tests 22

House rules for the code 22

10. Rules the code never breaks 22

11. Troubleshooting 22

12. Ethics and limits 24

Glossary 24

1. What is Sharko?

Sharko is a small program you run in the terminal. It answers one question about a loan application:

> **The question Sharko answers**
> If this loan application would be rejected, what is the smallest change to the loan amount or the loan term that would get it approved?

To answer it, Sharko learns from a public dataset of 4,269 past loan decisions (the Loan Approval Prediction Dataset from Kaggle, in an Indian lending setting). From those decisions it builds a model: a formula that gives any application a chance of approval from 0% to 100%. A chance of 50% or more counts as approved. All amounts are in Indian Rupees (INR).

A worked example

Here is a real run of the built-in example application:

Education	Graduate	Annual income	4,100,000 INR
Self-employed	No	CIBIL credit score	417
Dependents	2	Loan amount asked for	12,200,000 INR
Assets (home, business, luxury, bank)	2,700,000 / 2,200,000 / 8,800,000 / 3,300,000	Loan term asked for	8 years

What Sharko tried	Result
The application as submitted	0.3% chance of approval. Not approved.
Change the loan amount only	Nothing found. Sharko tried all 392 other amounts.
Change the loan term only	Nothing found. Sharko tried all 9 other terms.
Change both amount and term	Found. 14,600,000 INR over 4 years gives a 54.1% chance (approved).

What Sharko is not

- **Not a lender.** The result is a model-generated what-if, never a loan offer or a decision.
- **Not financial advice.** A change that works for the model is a pattern in past data, not a recommendation.
- **Not a guarantee.** A high chance of approval is the model's estimate. It is not a promise that a real lender would approve.
- **Not a form filer.** Sharko never submits anything. It only reads, scores and reports.

> **Why can a bigger loan become approved?**
> The model only follows patterns in the past loans it studied. A bigger loan can score higher because of those patterns, not because bigger loans are safer. Sharko says this in the result whenever the suggested amount is larger than the one asked for.

2. How Sharko works

Sharko works in four steps. Steps 1 and 2 always happen. Steps 3 and 4 happen only when the first answer is "not approved".

1 Learn
The model studies 4,269 past loan decisions. Run once.

2 Check
You enter one application. The model gives a chance of approval. Every time you check.

3 Search
Try other amounts and terms, closest to yours first.

4 Compare
Three ways side by side, with the closest marked.

Steps 3 and 4 run only when the chance of approval is below 50%.

1. Learn. The `train` command shows the model the past loans. Sharko tries three kinds of model (logistic regression, random forest, gradient boosting), tests each with a repeated practice test, and keeps the best one. In this project the winner is gradient boosting. You do this once.

2. Check. You answer 11 short questions about one application. The model returns a chance of approval. At 50% or more the answer is approved, the search is skipped, and you get a report showing how safe the approval is.

3. Search. If the chance is below 50%, Sharko tries other loan amounts and terms, nearest to your request first. It scores each option with the same model and stops at the first one that reaches 50%.

4. Compare. Sharko runs the search three ways and shows them side by side. The way with the smallest change is marked as the closest.

What stays fixed and what can change

Never changes during a search (fixed)	Can change during a search
Education, self-employed, number of dependents, annual income, CIBIL score, and the four asset values (residential, commercial, luxury, bank)	Loan amount and loan term only

The three ways of searching

Way	What changes	How many options it can try (real data)
Amount only	The loan amount. The term stays as submitted.	392: every amount from 300,000 to 39,500,000 in steps of 100,000, except yours
Term only	The loan term. The amount stays as submitted.	9: the terms 2, 4, 6 ... 20 years, except yours
Amount and term	Both. Each option must differ from yours in both.	Up to 3,528 (392 amounts x 9 terms)

How Sharko decides what is "closest"

Every option gets a distance from your request. A distance of 0 means no change; bigger means a bigger change. Sharko tries options from the smallest distance up, so the first approved option is also the smallest change.

> **distance**
> |new amount - old amount| / (largest amount - smallest amount) + |new term - old term| / (longest term - shortest term)

Worked example, from the run in section 1 (old: 12,200,000 over 8 years, new: 14,600,000 over 4 years):

Part	Calculation	Result
Amount part	2,400,000 / (39,500,000 - 300,000) = 2,400,000 / 39,200,000	0.0612
Term part	4 years / (20 - 2) = 4 / 18	0.2222
Distance	0.0612 + 0.2222	0.2834

> **Ties and order**
> If two options have the same distance, Sharko tries the smaller amount first, then the shorter term. Ties are decided on exact whole numbers, never on rounded decimals, so the same input always gives the same answer.

3. Quick start

Run every command from the Sharko project folder. You need Python 3.10 or newer.

0. Install (once)

```
python -m pip install -r requirements.txt
```

This installs scikit-learn, pandas, numpy, matplotlib, joblib and pytest.

1. See the start screen

```
python -m sharko
```

With no command, Sharko prints this short guide instead of an error:

```
======================================================================
SHARKO - loan what-if tool
======================================================================
Sharko predicts whether a loan application would be approved. If it
would be rejected, Sharko looks for the smallest change to the loan
amount or the loan term that would be approved.
The amounts used here are in Indian Rupees (INR).

QUICK START (2 steps)
----------------------------------------------------------------------
  1. python -m sharko train
       Teach Sharko from the loan data. Run this once.
  2. python -m sharko check --example
       See a full worked example, no typing needed. It also
       saves a report you can open: look in the Reports folder.

THEN
----------------------------------------------------------------------
  python -m sharko check           Enter your own application
  python -m sharko tutorial        Step-by-step guide
```

```
  python -m sharko history         Look back at earlier checks
  python -m sharko experiment      Test many loans (needed for charts)
  python -m sharko report --open   Open the charts in your browser

All commands: train, evaluate, check, history, experiment,
robustness, scenarios, report, tutorial. Add --help to any
command for details.
```

2. Teach Sharko (once)

```
python -m sharko train
```

Sharko learns from `loan_approval_dataset.csv` and saves the model to `artifacts/model.joblib`. Run it again only if the data changes.

3. Try the built-in example

```
python -m sharko check --example
```

No typing needed. Sharko checks a sample application, prints the result, and saves a report in the Reports folder.

4. Check your own application

```
python -m sharko check
```

Sharko asks for your first and last name (they only name your report file), then 11 short questions. Press Ctrl+C at any time to quit.

5. See the charts for all test loans (optional)

```
python -m sharko experiment
python -m sharko report --open
```

Run `experiment` first. It tests many loans and takes a short while. Then `report` draws four charts and opens them in your browser.

6. Check how steady and reliable the results are (optional)

```
python -m sharko scenarios
python -m sharko robustness
python -m sharko evaluate
```

These are for analysts. Section 7 explains what each one shows.

> **Something went wrong?**
> The most common messages and their fixes are in section 11, Troubleshooting.

How long things take

Measured on Windows 11 with Python 3.14. A slower machine, or the first run, takes longer. The install time depends on your connection.

Command	About
`python -m pip install -r requirements.txt`	1 to 3 minutes
`python -m sharko train`	10 seconds
`python -m sharko evaluate`	5 seconds
`python -m sharko check --example`	6 seconds
`python -m sharko experiment`	25 seconds
`python -m sharko scenarios`	25 seconds
`python -m sharko report --open`	6 seconds
`python -m sharko robustness`	3 minutes (6 full runs)
`python -m pytest`	2 to 3 minutes

4. Using the check command

`python -m sharko check` is the command you will use most. It scores one application, runs the search if needed, prints the result, and saves your report.

The questions

Sharko first asks for your name, then these 11 questions. Each question shows its allowed answers. A wrong answer is explained and asked again.

#	Question	Allowed answers
-	First name, Last name	Letters, spaces, hyphens and apostrophes, 1 to 40 characters. Only used to name your report file.
1	Education	1 = Graduate, 2 = Not Graduate (or type the name)
2	Self-employed	1 = Yes, 2 = No (also y or n)
3	Number of dependents	A whole number, 0 or more
4	Annual income (INR)	A whole number, 0 or more. Commas are allowed, for example 4,100,000
5	CIBIL credit score	A whole number from 300 to 900
6	Residential assets value (INR)	A whole number, 0 or more
7	Commercial assets value (INR)	A whole number, 0 or more
8	Luxury assets value (INR)	A whole number, 0 or more
9	Bank asset value (INR)	A whole number, 0 or more
10	Loan amount (INR)	300,000 to 39,500,000 in steps of 100,000
11	Loan term (years)	One of 2, 4, 6, 8, 10, 12, 14, 16, 18, 20

> **Where do the loan limits come from?**
> From the training data. Sharko prints them at the start when it asks the questions. They are also the limits of the search, so Sharko only suggests amounts and terms that appear in real past loans.

Skipping the questions with flags

Every question has a flag with the same name. Give all 11 with `--no-prompt` and Sharko asks nothing:

```
python -m sharko check --no-prompt --first-name Ana --last-name Cruz --education Graduate
  --self-employed No --no-of-dependents 2 --income-annum 4100000 --cibil-score 417
  --residential-assets-value 2700000 --commercial-assets-value 2200000
  --luxury-assets-value 8800000 --bank-asset-value 3300000 --loan-amount 12200000
  --loan-term 8
```

Flag	What it does
`--example`	Use the built-in sample application. Any flag you add overrides one value, for example `--cibil-score 800`.
`--no-prompt`	Never ask questions. Every answer must be a flag. Missing names become Unnamed Applicant.
`--first-name`, `--last-name`	Your name, for the report title and file name.
`--open`	Open the saved report in your browser without asking.
`--reports-dir FOLDER`	Where to save the report. Default: Reports
`--out FOLDER`	Where the experiment results are, so your report can include the paper's charts. Default: outputs
`--model FILE`	The saved model. Default: artifacts/model.joblib
`--history FILE`	The log of checks. Default: artifacts/history.jsonl
`--education`, `--self-employed`, `--no-of-dependents`, `--income-annum`, `--cibil-score`, `--residential-assets-value`, `--commercial-assets-value`, `--luxury-assets-value`, `--bank-asset-value`, `--loan-amount`, `--loan-term`	One flag per question, in the same order as the table above.

Reading the result in the terminal

Here is a real result (the built-in example). Read it from the top; you can stop at IMPORTANT.

```
RESULT: Not approved as submitted
------------------------------------------------------------------------
  The model gives this application a 0.3% chance of approval. It needs
  50.0% or more to be approved.

SMALLEST CHANGE THAT WOULD BE APPROVED
------------------------------------------------------------------------
  Change both the amount and the term:
    Loan amount      12,200,000 -> 14,600,000 INR
    Loan term        8 -> 4 years
    Approval chance  0.3% -> 54.1% (approved)
  Note: a bigger loan is not a safer loan. The model only follows
  patterns in past loans, so treat this as a pattern, not as advice.

HOW SHARKO FOUND THIS
------------------------------------------------------------------------
  1. Kept everything else about the applicant the same. Only the loan
     amount and the loan term were allowed to change.
  2. Listed the allowed amount and term options, closest to the original
     request first.
  3. Asked the model to score each option, one by one. Sharko stops at
     the first option with a 50.0% chance or more. For the closest way,
     that was try 363.
  4. Did this 3 ways (amount only, term only, both) and compared them.

THE 3 WAYS COMPARED
```

```
------------------------------------------------------------------------
  Way              Result   Options tried
  Amount only      none               392
  Term only        none                 9
  Amount and term  found              363  <- closest

IMPORTANT
------------------------------------------------------------------------
Model-generated what-if result. This is not a lender decision,
a loan offer, or financial advice. A high chance of approval is
an estimate, not a guarantee that a real lender would approve.
Limitations: the model reflects patterns in a public dataset that
lacks economic variables (e.g. inflation, interest rates) and
demographic detail. If no approved option is found, that means none
was found within the searched amounts and terms, not that none exists.
Scores can change with a different data split; run
python -m sharko robustness to see by how much.
```

Section	What it tells you
RESULT	Approved or not, with the model's chance of approval and the 50% cut-off.
SMALLEST CHANGE THAT WOULD BE APPROVED	The closest amount and term that work, with the chance before and after.
HOW SHARKO FOUND THIS	The four steps behind the answer, with this run's numbers.
THE 3 WAYS COMPARED	Amount only, term only, and both, with how many options each tried. The closest way is marked.
IMPORTANT	The disclaimer and the limits of the model. Always shown.
TECHNICAL DETAILS (for analysts)	Exact scores, distances, search limits, model quality and a short glossary. Safe to skip.
REPORT SAVED	Where your HTML report was written.

If the application is already approved

Sharko skips the search and says so. It also reminds you that the percentage is an estimate. This is a real run of `python -m sharko check --example --cibil-score 800 --loan-amount 5000000 --loan-term 12`:

```
RESULT: Approved as submitted
------------------------------------------------------------------------
  Congratulations! The model gives this application a 99.8% chance of
  approval (it needs 50.0% or more). No change to the loan is needed.
  This is the model's estimate, not a promise that a real lender will
  approve it.

HOW THIS WAS SCORED
------------------------------------------------------------------------
  The model compares your answers with patterns in past loan decisions
  and gives a chance of approval. Because the answer is already
  approved, Sharko did not search for changes.
```

The end of every check

```
REPORT SAVED
------------------------------------------------------------------------
  Your report has your answers, the result, and every try the search
  made, step by step. It also has the paper's charts for all the test
  loans.
  Folder: C:\Users\you\Sharko\Reports
  File:   Example_Applicant_Results.html
  Open it by double-clicking the file, or in any web browser.
  Tip: add --open to open it automatically next time.
```

5. Your report file

Every check saves an HTML report by itself, so there is nothing extra to run. It is one file with everything inside, so you can open it in any browser or send it to a teammate.

Where it is saved and what it is called

- **Folder:** Reports, inside the Sharko folder. The terminal prints the full path.
- **File name:** `FirstName_LastName_Results.html`, for example `Ana_Cruz_Results.html`.
- Spaces become underscores. Accents and symbols are removed, so a name like José becomes Jose.
- Running check again with the same name saves `..._Results_2.html`, then `_3`, so nothing is overwritten.
- Open it by double-clicking the file, answer `y` when Sharko asks, or add `--open`.

What is inside

Section	Not approved	Approved
What you submitted	All 11 answers, marked Can change or Fixed	Same
The result	Chance, the 50% cut-off, the smallest change (before and after)	Chance and a congratulation, with the note that it is an estimate
How Sharko found this	The 4 steps with this run's numbers	How this was scored
The 3 ways compared	Result, options tried, chance after, distance	-
Try by try	Every option the search scored: first tries, last try, and a chart	-
How safe is this approval?	-	Two charts: what happens if only the term, or only the amount, changes
How this compares with the experiment	The paper's four charts and a you-versus-the-experiment table	The paper's four charts
Model quality, Important, Words used	Scores, disclaimer and limits, glossary	Same

Try by try chart (not approved). One panel per way. The dashed line is the 50% needed. A star marks the first option that reached it, where Sharko stopped. The labels stay inside their own panel.

Safety-margin chart (approved). Shows how much room an approval has. The orange dot is your application; the grey area below the dashed line is not approved.

> **Privacy**
> Reports contain your name and financial details, so the Reports folder is git-ignored and never uploaded. Names are used only to title and name the report. They are not given to the model and are not written to the history log. A sample report with a made-up applicant is kept in `samples/reports`.

6. The experiment and the charts

The check command looks at one application. The experiment looks at many. It takes the 854 loans that were set aside for testing (never used for learning), keeps the 312 that the data marks Rejected and the model also rejects (the "eligible" loans), and runs all three searches on every one of them.

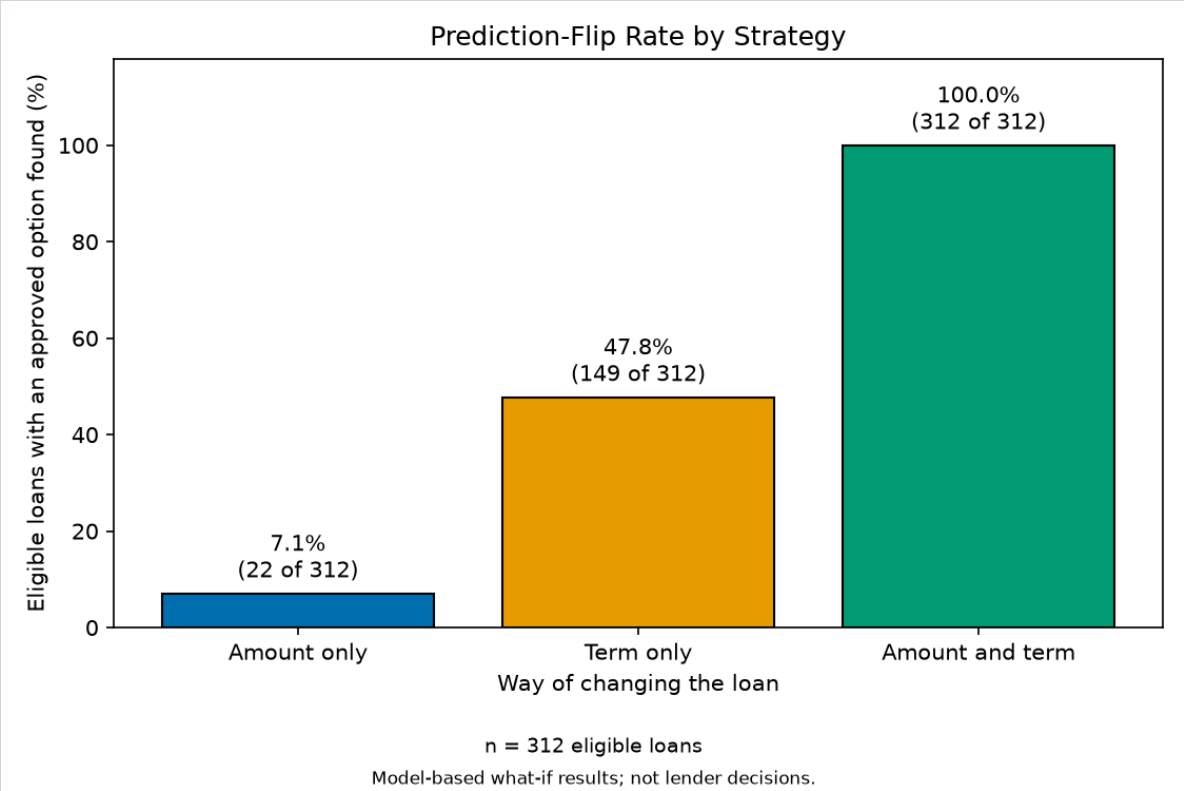
Way	Found an approved option	Median options tried	Median distance (when found)
Amount only	7.1% (22 of 312)	392	0.1110
Term only	47.8% (149 of 312)	9	0.4444
Amount and term	100.0% (312 of 312)	1,053	0.4795

Results from the saved experiment on the real data. They describe this trained model, not lenders in general.

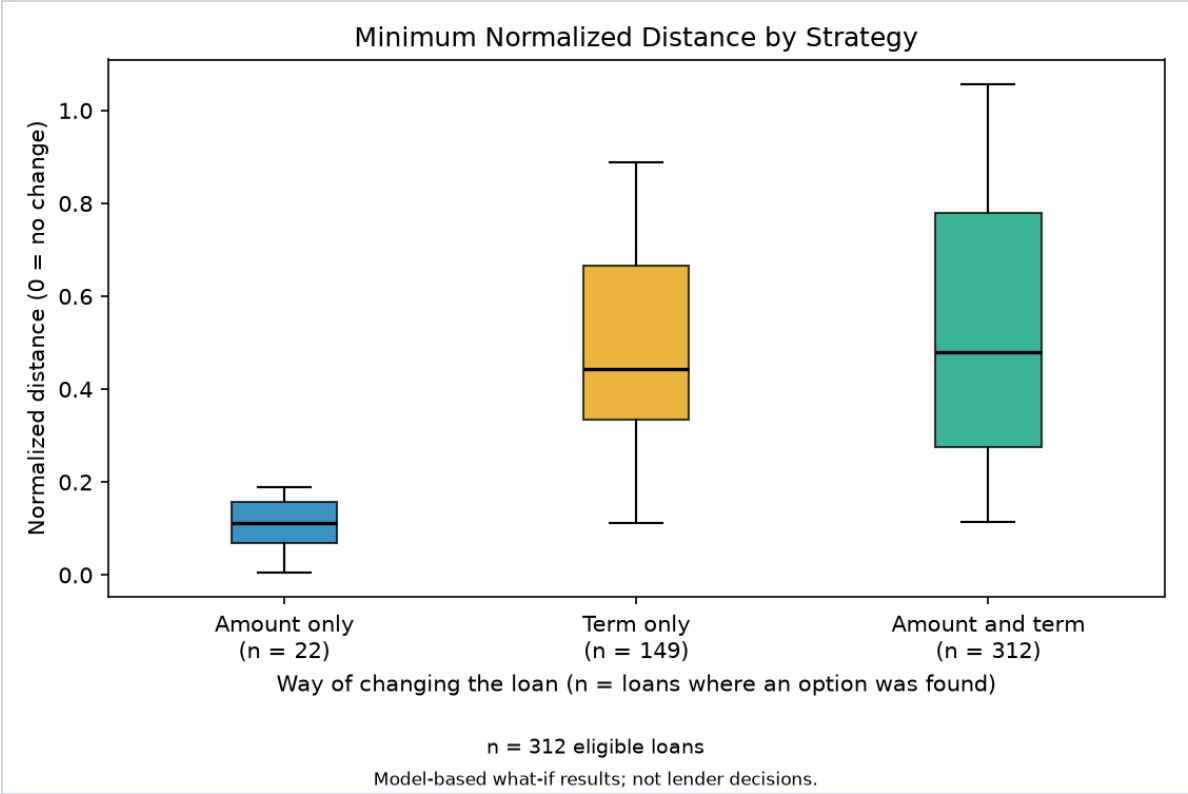
```
python -m sharko experiment
python -m sharko report --open
```

`experiment` writes `outputs/experiment_results.csv`. `report` turns it into tables and the four charts, and opens `Reports/Experiment_Report.html`, one page with all four charts and a plain caption for each.

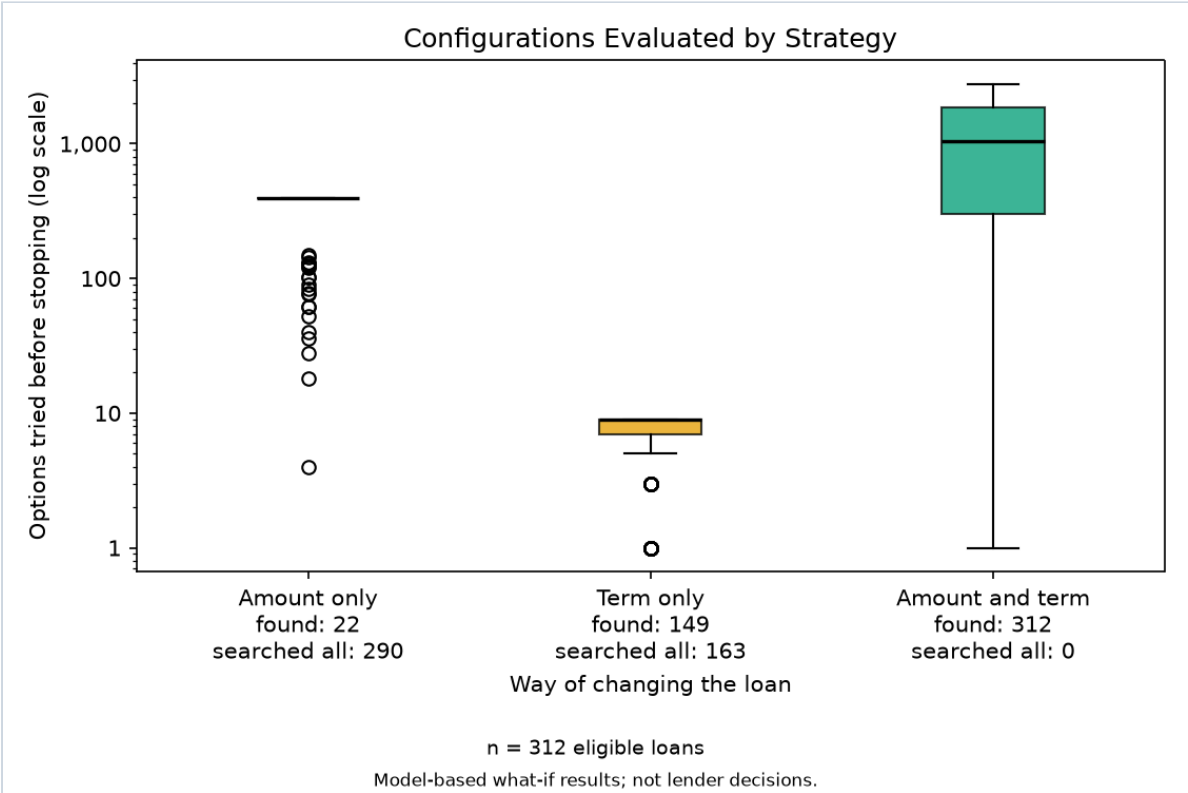
The four charts and how to read them



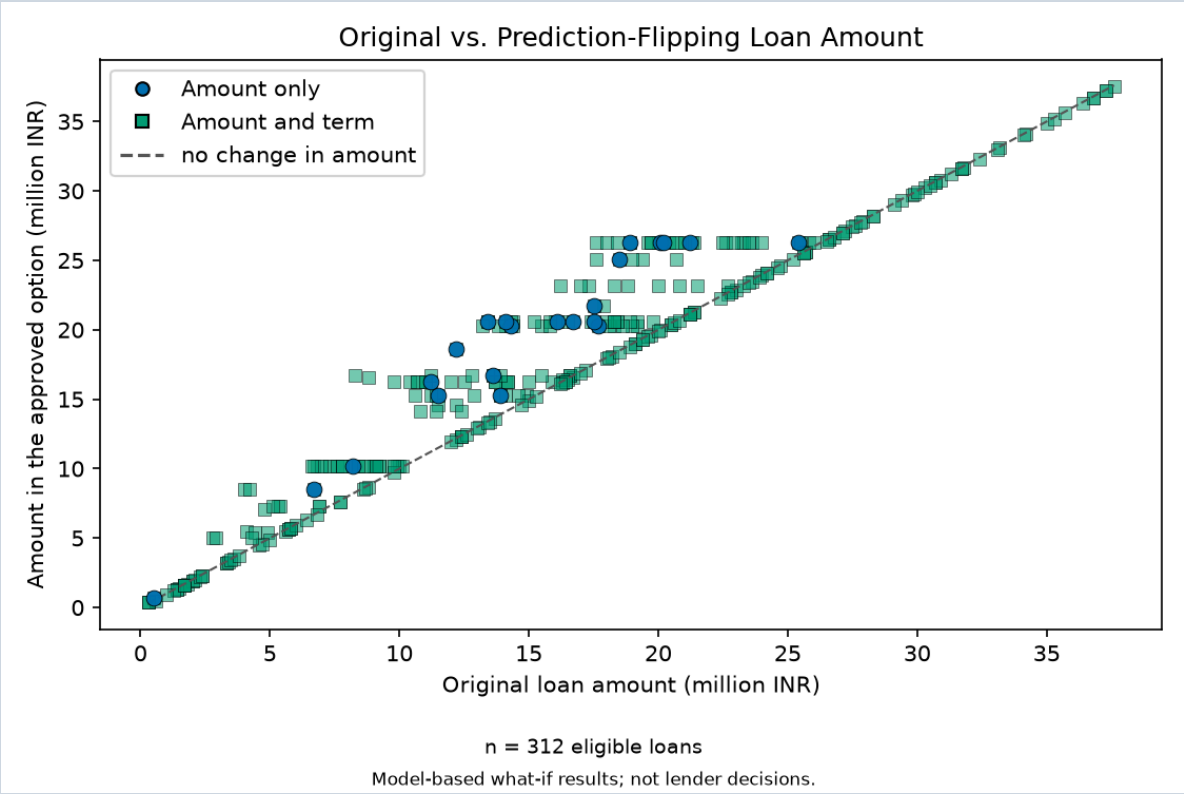
Prediction-flip rate by strategy. The share of eligible loans where a way found an approved option. A taller bar means that way worked more often.



Minimum normalized distance by strategy. How big the needed change was, only for loans where a change was found. A lower box means a smaller change. The n under each box is how many loans it covers.



Configurations evaluated by strategy. How many options were tried before stopping, on a log scale. "Found" loans stopped early; "searched all" loans tried every option without success.



Original vs prediction-flipping loan amount. Each dot is a loan whose search changed the amount. A dot on the dashed line means the amount did not change.

Files the experiment and report write

File	What it holds
`outputs/experiment_results.csv`	One row per eligible loan and way: the option found, distance, score change, options tried
`outputs/summary.csv`	One row per way with all the averages
`outputs/paired_comparison.csv`	Ways compared on the loans that both could fix
`outputs/*.png`	The four charts
`Reports/Experiment_Report.html`	All four charts on one page, with captions

7. How steady and how reliable are the results?

A model can look good once and still be fragile. Three analyst commands test that. Every number below comes from a real run on the full dataset, and the numbers are the same each time because the whole pipeline uses a fixed seed (42).

Negative assets and different random splits

The dataset has 28 loans with a negative residential asset value, which is not possible for a real asset. `python -m sharko robustness` asks two questions:

- **Do those 28 loans matter?** It trains, tests and runs all three searches twice: once on every loan, and once without those 28.
- **Does the result depend on one lucky split?** It repeats the whole run with 5 different random splits (seeds 42 to 46). Use `--seeds 10` for more.

Run	Loans	Accuracy	ROC-AUC	Amount only	Term only	Both
All records	4,269	0.9824	0.9983	7.1%	47.8%	100%
Without negative assets	4,241	0.9823	0.9979	8.9%	48.9%	100%

Seed (split)	Accuracy	ROC-AUC	Amount only	Term only	Both
42	0.9824	0.9983	7.1%	47.8%	100%
43	0.9789	0.9970	9.7%	54.0%	100%
44	0.9848	0.9985	8.9%	47.0%	100%
45	0.9719	0.9939	8.4%	50.0%	100%
46	0.9836	0.9982	6.3%	53.2%	100%
Mean (std)	0.9803 (0.005)	-	8.1% (1.4)	50.4% (3.2)	100% (0)

The last three columns are the share of eligible test loans where that way found an approved option.

> **What this means**
>
> Removing the 28 loans changes almost nothing (accuracy 0.982 in both runs, and the three flip rates move by about 2 points at most). Across five splits the accuracy stays between 0.972 and 0.985 and the order of the three ways never changes: both is best, then term only, then amount only. The results do not depend on those 28 loans or on one lucky split. The two files are saved as `outputs/robustness_sensitivity.csv` and `outputs/robustness_seeds.csv`.

Different kinds of applicant

`python -m sharko scenarios` compares the three ways for different kinds of applicant. It uses the same 312 eligible test loans as the experiment and prints four worked examples picked from them:

Example	Asked for	Chance as asked	Amount only	Term only	Amount and term
Lowest CIBIL score (300)	17,000,000 over 8 years	0.1%	none found	none found	23,200,000 over 4 years (52.0%)
Middle CIBIL score (434)	18,300,000 over 8 years	0.2%	none found	none found	20,600,000 over 4 years (52.1%)
Highest CIBIL score (548)	13,700,000 over 10 years	1.7%	none found	13,700,000 over 4 years (90.4%)	13,600,000 over 4 years (90.4%)
Largest loan for the income (CIBIL 507)	6,000,000 over 6 years	0.4%	none found	6,000,000 over 4 years (73.2%)	5,900,000 over 4 years (73.2%)

It then splits the 312 loans into the lowest, middle and highest third of CIBIL score, annual income and loan amount relative to income, and shows how often each way found an approved option (104 loans per group):

Grouped by	Group (range)	Amount only	Term only	Both
CIBIL credit score	Lowest third (300-382)	5.8%	48.1%	100%
	Middle third (383-472)	8.7%	48.1%	100%
	Highest third (473-548)	6.7%	47.1%	100%
Annual income	Lowest third (0.2M-4.1M)	3.8%	52.9%	100%
	Middle third (4.1M-7.0M)	7.7%	45.2%	100%
	Highest third (7.0M-9.9M)	9.6%	45.2%	100%
Loan amount / annual income	Lowest third (1.50-2.47)	15.4%	13.5%	100%
	Middle third (2.47-3.34)	5.8%	37.5%	100%
	Highest third (3.34-4.00)	0.0%	92.3%	100%

> **What this means**
>
> The CIBIL score barely changes how well each way works, and income changes it a little. The loan size relative to income changes it a lot. Changing only the term works for just 13.5% of the smallest loans for the income but 92.3% of the largest. Changing only the amount works best for the smallest loans (15.4%) and never for the largest (0.0%). Changing both found an option for every loan in every group. The full table, with distances, is saved as `outputs/scenarios.csv`. These are patterns in this model and dataset, not advice.

Do the percentages mean what they say?

A model can rank loans well and still give percentages that are too high or too low. `python -m sharko evaluate` checks this. It prints the Brier score (the average squared gap between the model's chance and what happened; 0 is perfect, and always answering 50% scores 0.25) and groups the 854 test loans by the chance the model gave them:

Chance given	Loans	Model said	Truly approved
0% to 10%	292	0.4%	0.0%
10% to 20%	10	13.8%	0.0%

Chance given	Loans	Model said	Truly approved
20% to 30%	4	25.2%	25.0%
30% to 40%	7	34.8%	14.3%
40% to 50%	3	46.1%	66.7%
50% to 60%	2	53.6%	50.0%
60% to 70%	6	65.2%	100.0%
70% to 80%	13	76.1%	76.9%
80% to 90%	9	85.0%	44.4%
90% to 100%	508	99.5%	99.6%

What this means

The Brier score is 0.0144 and the average gap is 1.3 percentage points, so the percentages are reliable at the ends: loans scored under 10% or over 90% were almost always rejected or approved as predicted. The middle is thin. Only 54 of the 854 test loans scored between 10% and 90%, and the 80% to 90% group was truly approved less than half the time (9 loans). Sharko stops at the first option with 50% or more, so a suggested change usually lands in that middle zone. Read a suggested "52% chance" as a weaker signal than a "99% chance". Small groups are noisy, so do not read much into any one row.

Sample outputs

A copy of one full run is kept in the `samples` folder, so you can see every result without running anything. It holds the terminal output of each command, all the CSV files, the four charts, the experiment page and a report for the built-in example. `samples/README.md` lists each file and the commands that make it.

8. Command reference

Every command starts with `python -m sharko`. Add `--help` to any command for its options.

Command	What it does	Main options (defaults)
(none)	Prints the quick-start screen.	-
`tutorial`	Prints a step-by-step guide to installing, running and reading Sharko.	-
`train`	Tries three kinds of model, keeps the best, and saves it. Prints practice-test scores and final test scores.	`--data` (loan_approval_dataset.csv), `--model` (artifacts/model.joblib)
`evaluate`	Re-tests the saved model on the held-back loans. Prints the scores, the Brier score and the calibration table.	`--data`, `--model`
`check`	Scores one application, searches if needed, prints the result, saves the report and logs the run.	`--example`, `--no-prompt`, `--first-name`, `--last-name`, `--open`, `--reports-dir` (Reports), `--out` (outputs), `--model`, `--history`, plus one flag per question
`history`	Lists your earlier checks: time, result, chance, and the closest fix.	`--history` (artifacts/history.jsonl), `--last` (10)
`experiment`	Runs all three searches on every eligible test loan and saves the results.	`--data`, `--model`, `--out` (outputs)
`scenarios`	Shows four worked examples and compares the three ways by CIBIL score, income and loan size. Saves `scenarios.csv`.	`--data`, `--model`, `--out` (outputs)
`robustness`	Compares the results with and without the 28 negative-asset loans, and across random splits. Saves two CSV files.	`--data`, `--seeds` (5), `--out` (outputs)
`report`	Builds tables, the four charts and the experiment page from the experiment results.	`--out` (outputs), `--reports-dir` (Reports), `--open`

What the scores printed by train and evaluate mean

```
Testing the saved model on the 20% of loans it never saw...
  accuracy  0.9824  share of loans predicted correctly
  precision 0.9796  when it says Approved, how often it is right
  recall    0.9925  share of truly Approved loans it caught
  F1        0.9860  balance of precision and recall
  ROC-AUC   0.9983  ranks Approved above Rejected (1 = perfect)
  Brier     0.0144  Brier score, lower is better (0 = perfect)
  Hits and misses on the test loans:
    Truly Rejected: 312 right, 11 wrong
    Truly Approved: 527 right, 4 wrong
```

These come from the 20% of loans the model never saw while learning. ROC-AUC near 1 means the model ranks approved loans above rejected ones almost perfectly. The Brier score measures how close the

percentages are to what happened. Section 7 explains the calibration table that follows these scores in the `evaluate` output.

Exit codes

Code	Meaning
0	Finished normally.
2	Something to fix: a missing model or dataset, or an answer that is not allowed. The message says what.
130	You pressed Ctrl+C or the input ended. Nothing was saved.

9. How the code is organized

The code lives in the `sharko` folder. Each file does one job. This map shows how a command flows through them:

1 Start

cli.py
Reads your command, then runs the steps below in order. Every command starts here: python -m sharko ...

2 Do the work

intake.py
Asks the questions and checks the answers

model.py
The trained model: chance of approval and calibration

search.py
The what-if search and the distance

trace.py
Try-by-try list and safety margin

3 Show and save

render.py
The result you read in the terminal

applicant_report.py
Your report file (HTML)

history.py
A log of your checks

report.py
Experiment charts (with dashboard.py)

4 Analyst checks

experiment.py
Runs the search on many loans

scenarios.py
Compares the ways by applicant profile

robustness.py
Repeats the run without odd records and with other seeds

File by file

File	Its job	Main pieces
`__main__.py`	Lets you run the app with `python -m sharko`.	-
`cli.py`	Defines every command and runs the steps in order.	build_parser, main, one function per command
`config.py`	Fixed settings in one place.	SEED 42, APPROVAL_THRESHOLD 0.5, AMOUNT_STEP 100,000, FEATURES, default paths
`data.py`	Loads the dataset, splits it 80/20, works out the search limits.	load_dataset, split_dataset, SearchSpace
`model.py`	Trains and saves the model, scores applications, and checks how reliable the percentages are.	train_model, evaluate_model, approval_scores, predicted_status, calibration_table, calibration_error

File	Its job	Main pieces
`search.py`	The what-if search and the distance.	Strategy, generate_candidates, normalized_distance, run_all_strategies, SearchResult
`experiment.py`	Runs the search on many loans and summarizes.	eligible_applications, run_experiment, summarize, paired_comparison
`scenarios.py`	Splits the experiment results by applicant profile and picks the worked examples.	profile_breakdown, pick_examples, PROFILE_DIMENSIONS
`robustness.py`	Runs the whole pipeline with and without the negative-asset loans and over several seeds.	without_negative_assets, run_once, sensitivity, seed_check, spread
`intake.py`	Asks the questions and checks every answer.	validate_field, collect_applicant, collect_name, report_filename
`trace.py`	Rebuilds the try-by-try list and the safety margin for reports.	build_trace, margin_scan
`render.py`	Writes the result you read in the terminal.	format_check_result
`plain.py`	Plain-language wording helpers.	chance, wrap, search_steps, PLAIN_WAYS
`messages.py`	The disclaimer, limits note and error messages.	DISCLAIMER, LIMITATIONS_NOTE
`guide.py`	The start screen and the tutorial text.	start_screen, tutorial
`applicant_report.py`	Builds your report file.	write_applicant_report, ExperimentContext
`report.py`, `dashboard.py`	Draw the four charts and the experiment page.	generate_report, render_figures, write_dashboard
`history.py`	Appends each check to a log file. Nothing ever reads it back to change a result.	HistoryLog

Where do I change...?

I want to change	Look in
The words printed after a check	render.py and plain.py
The disclaimer or the limits note	messages.py
The start screen or tutorial	guide.py
A question, its allowed answers or its error text	intake.py
The search limits or the split	data.py and config.py
How a candidate is ordered or the distance	search.py (changes the results, so run the tests)
The models tried or how the best is chosen	model.py
The calibration table or the Brier score	model.py (and the printing in cli.py)
The applicant groups in `scenarios`	scenarios.py

I want to change	Look in
The checks in `robustness`	robustness.py
The look of your report page	applicant_report.py
The experiment charts	report.py

Tests

Run all tests with `python -m pytest`. There are 215, and they take 2 to 3 minutes. The project is built test first: a failing test is written, then the code that makes it pass. Tests run inside a scratch folder, so they never leave files in your project.

House rules for the code

- Run everything as `python -m ...`, with no shell-specific syntax.
- Use pathlib for paths and pass `encoding="utf-8"` to every text file.
- Anything printed in the terminal is plain ASCII, so it looks the same on every system.
- Terminal output stays within 72 columns so it never wraps badly.

10. Rules the code never breaks

These rules come from the project paper. They keep the results fair, repeatable and honest.

- Only the loan amount and loan term change during a search. Every other field stays as submitted.
- The split is 80/20, stratified on the loan status, with seed 42. The search limits, the scaling and the choice of model use the 80% training part only. (`robustness` tries other seeds only to test how steady the results are. The official run stays on seed 42.)
- Approved means a model probability of 0.5 or more. The same model and threshold are used for every option.
- Options are tried in order of distance, then amount, then term. Ties are decided on exact whole numbers, never on decimals.
- No single "best" way is declared. Sharko shows the three ways side by side and marks the closest; it never combines them into one weighted score.
- Every result carries the disclaimer and shows the original and new amount and term, the model score and the distance.
- History is write-only. Nothing ever reads it back to change a search or the model.
- Sharko never submits anything and never changes the applicant's own details.

11. Troubleshooting

You see	What it means	What to do
Model not found at ...	The model has not been trained yet.	Run `python -m sharko train`

You see	What it means	What to do
Dataset not found at ...	Sharko is not running in the project folder.	Open the terminal in the Sharko folder, or add `--data PATH`
... must be one of / must be a whole number ...	An answer is outside the allowed list.	Use a value from the list shown under the question. Amounts are in steps of 100,000 and terms are even years from 2 to 20.
Experiment results not found	`report` needs the experiment first.	Run `python -m sharko experiment`, then `report`
Tip: run python -m sharko experiment ...	Your report was saved without the paper's charts.	Run the experiment, then run the check again
No module named 'sklearn' (or pandas, matplotlib)	The packages are not installed for this Python.	Run `python -m pip install -r requirements.txt` with the same `python` you use for Sharko
`python` is too old or not found	Sharko needs Python 3.10 or newer.	Check with `python --version`. On macOS or Linux try `python3`. On Windows try `py` in place of `python`.
The browser did not open	Your system blocked it.	Open the file in the Reports folder yourself
No history yet	No check has been run.	Run `python -m sharko check --example`
Cancelled. Nothing was saved.	You pressed Ctrl+C.	Run the command again when ready
`experiment`, `scenarios` or `robustness` seems stuck	It is working.	`experiment` prints progress. `scenarios` takes under a minute and `robustness` takes a few minutes.
Numbers differ slightly from this guide	A different library version, or a different data file.	Compare with `samples/README.md`, which lists the versions used

12. Ethics and limits

> **Read this before using any result**
>
> Model-generated what-if result. This is not a lender decision, a loan offer, or financial advice. A high chance of approval is an estimate, not a guarantee that a real lender would approve. The model reflects patterns in a public dataset that lacks economic variables (such as inflation and interest rates) and demographic detail. If no approved option is found, that means none was found within the searched amounts and terms, not that none exists. Scores can change with a different data split.

- The results describe how this trained model behaves. A different model or a real lender could answer differently.
- Sharko does not know if an applicant would accept a changed loan. A suggestion is never an offer.
- The dataset has 28 loans with a negative residential asset value. They are kept as they are. The robustness check (section 7) shows that removing them changes accuracy by 0.0001 and each flip rate by about 2 points at most.
- The results are steady across five random splits, but they come from one dataset. They should not be generalized automatically to other lenders or populations.
- The percentages are most reliable near 0% and 100% and least reliable in the middle, which is where a suggested change usually lands (section 7).
- Because the dataset has little demographic detail, Sharko makes no claim that the model is fair across groups.
- Historical lending patterns can carry their own biases, and the model learns them.

Glossary

Word	Plain meaning
Chance of approval	The model's estimate from 0% to 100% (also called the approval score, 0 to 1). 50% or more counts as approved.
Way (strategy)	One of the three searches: amount only, term only, or both.
Flip	A rejected application that becomes approved after a change. The "flip rate" is how often a way finds one.
Option, try, configuration	One amount-and-term pair the search scores. "Options tried" and "configurations evaluated" mean the same thing.
Distance (normalized distance)	How far an option is from your request. 0 means no change.
Eligible loan	A held-back test loan that the data marks Rejected and the model also rejects.
Held-out (test) data	20% of the loans, kept aside and never used for learning, so the scores are honest.
Cross-validation (CV AUC)	A repeated practice test on the training loans, used to choose the best model.
Accuracy, precision, recall, F1, ROC-AUC	Scores for how well the model predicts loans it never saw. Closer to 1 is better.
Brier score	The average squared gap between the model's chance and what happened. 0 is perfect; always answering 50% scores 0.25.
Calibration	Whether a "70% chance" is truly approved about 70% of the time.

Word	Plain meaning
Seed / split	A number that fixes how the loans are shuffled into the 80% and 20% parts. A different seed gives a different split.
Profile	A kind of applicant, such as the lowest third by CIBIL score.
CIBIL score	A credit score from 300 to 900.
INR	Indian Rupees.